

Term project #1: Matrix-vector multiplication w/ half-precision input data

05/10/2018

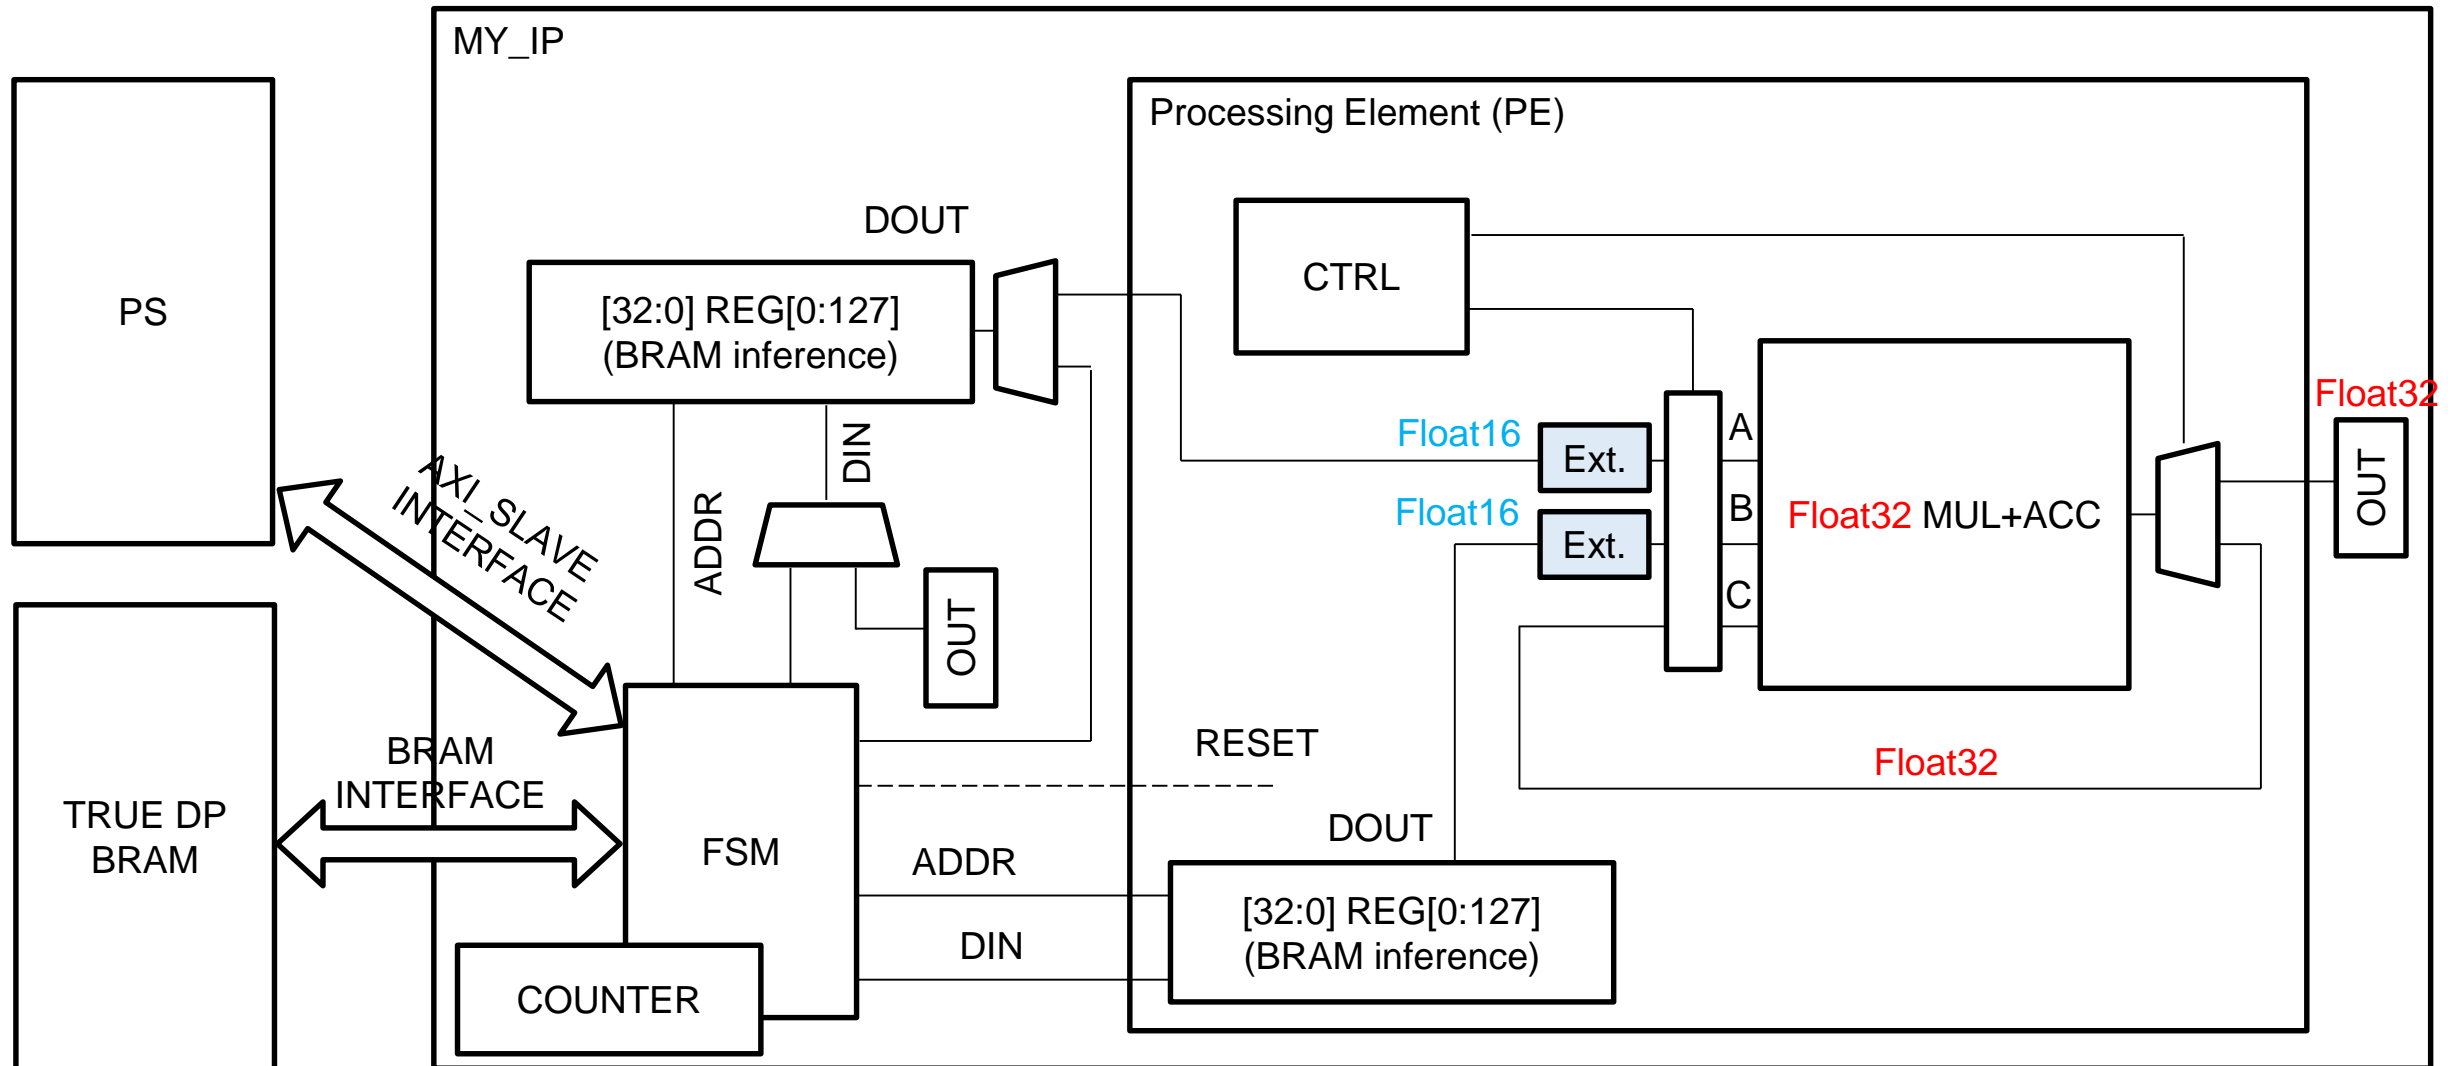
4190.309A: Hardware System Design
(Spring 2018)

Introduction to term project #1

- Matrix and vector multiplication
 - Multiply matrix and vector to produce vector
 - 4x4 MV multiplier example:

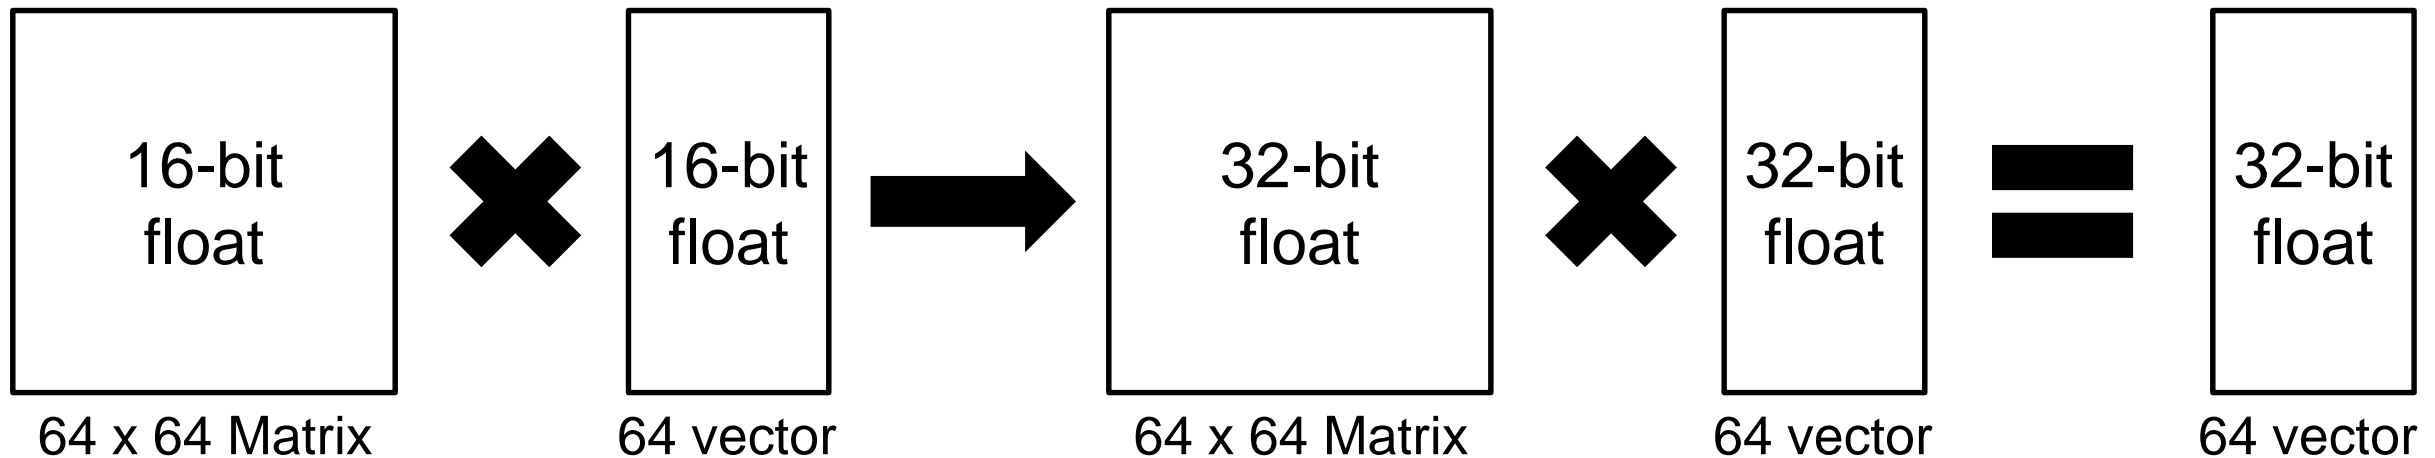
$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} \times \begin{bmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{bmatrix} = \begin{bmatrix} c_1 \\ c_2 \\ c_3 \\ c_4 \end{bmatrix}$$

Overview: Matrix-vector multiplication IP



Specification of Term Project #1

- **Goal:** Implement 64x64 MV multiplication IP with matrix and vectors described in *half-precision data format*
 - Input / output format is provided (See the next page)
 - Convert FLOAT16 into FLOAT32 (See p3, Ext. logic)
 - Calculate 64x64 matrix – 64 vector multiplication.



Specification of Term Project #1

■ **Scoring policy**

- *No scoring for performance now*, but we will perform an optional project
 - e.g. Performance tuning with a fixed point format
- Scoring of IP will be performed as follows
 - By precision lost of conversion, error less than 1% will be accepted as correct answer.
 - Penalties for incorrect results (e.g. missing partial results, ...)
 - Penalties for non-general functionality (e.g. work only once, ...)
 - Show your demo results in front of TAs
 - You must show tests at least 3 times with randomly generated inputs
 - **Due date: Lab time at 5/31 Thr. and 6/1 Fri.**


Specification of Term Project #1

- File description for software parts
 - `input.txt`
 - `generator.c` → `input_generator` (bin)
 - `main.cpp`, `zynq.cpp`, `zynq.h` → `project1` (bin)
 - `Makefile`
- File description for hardware parts (verilog)
 - Start from Lab 8
 - Design & edit IP implementation for 64x64 MV multiplication (`pe_controller.v`, etc.)
 - Generate bitstream
 - Execute *project1*


Specification of Term Project #1

- Data format

- **FLOAT32:** IEEE754 single-precision (**used for computation**)
 - 32bit (1bit sign / 8bit exponent / 23bit fraction)



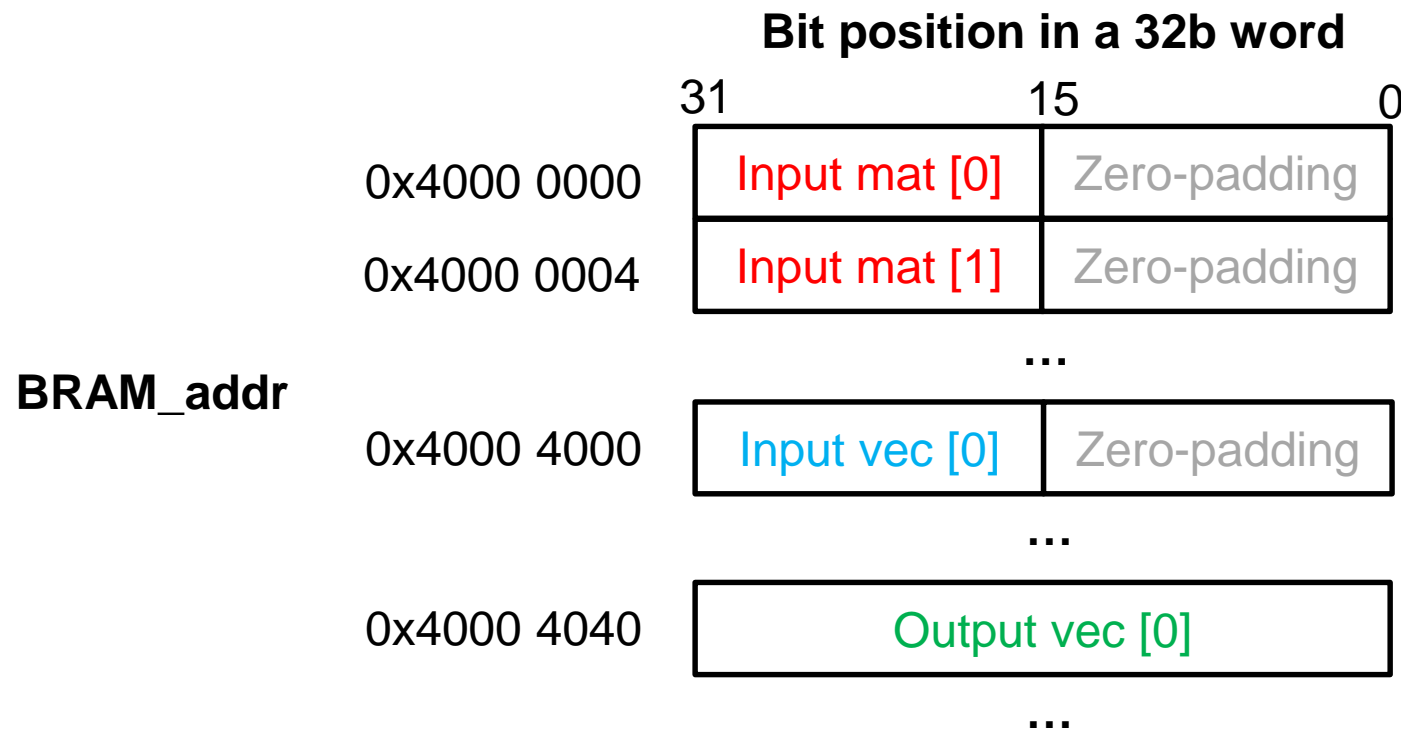
- **FLOAT16:** IEEE754 half-precision (**used for BRAM**)
 - 16bit (1bit sign / 5bit exponent / 10bit fraction)
 - We'll use this format, but have to convert it into FLOAT32 when computing with floating-point MAC IP



You must implement
FLOAT16 to FLOAT32
converter for internal use
in MV multiplication IP

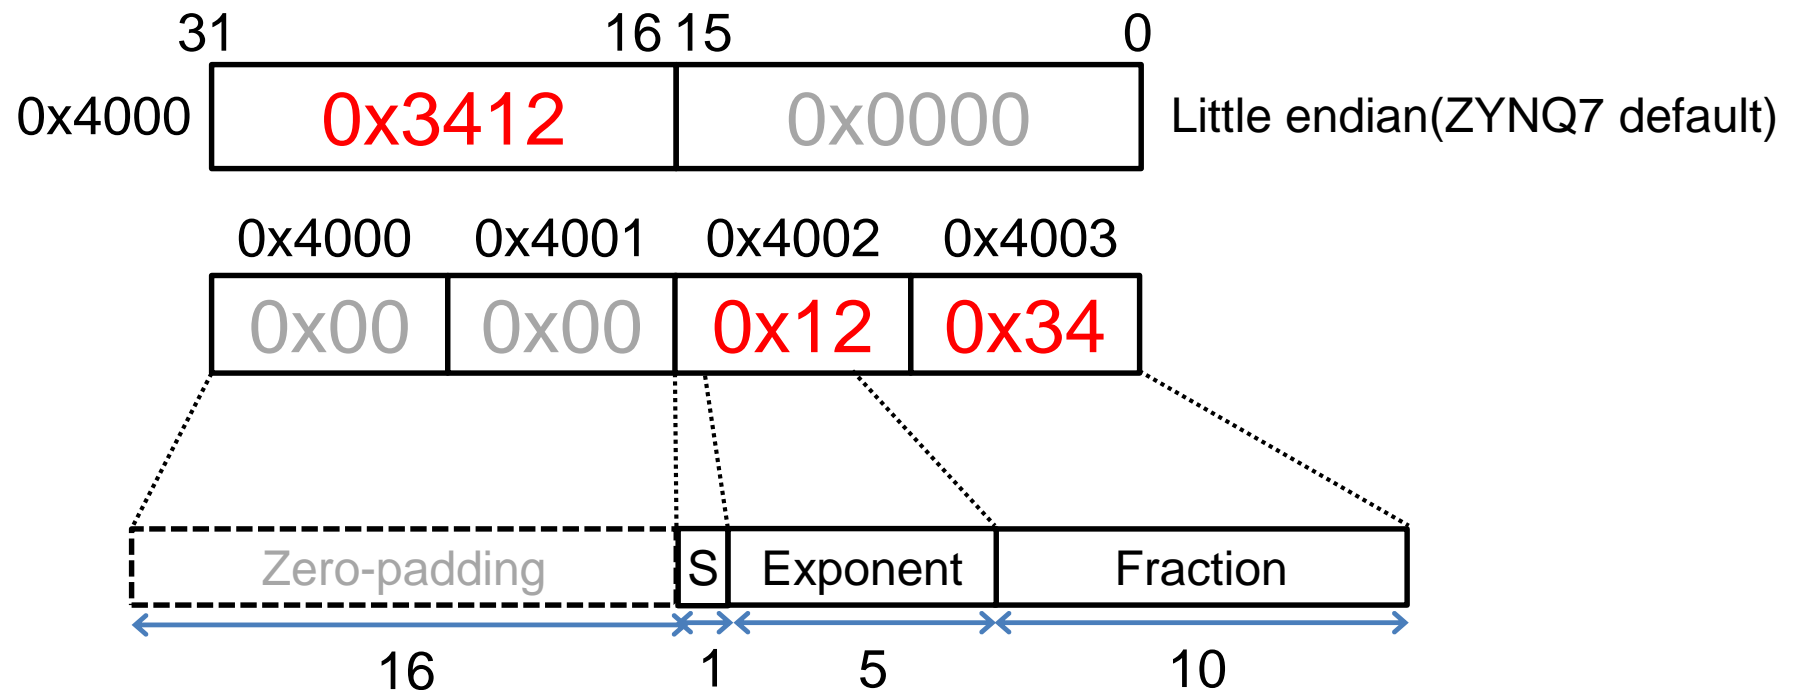
Specification of Term Project #1

- Input / output format
 - Input : 4 byte-aligned mini-float (16-bit data)
 - Valid data is stored at MSB [31:16]
 - Output should hold 32-bit floating point (Little-endian)



Specification of Term Project #1

- Closer look on input format
 - Example value : 0x1234
 - Looks like upper figure in C code
 - Actual byte order is like in second figure.



Specification of Term Project #1

- CPU-implemented code.
 - arm_calculate() function do the same thing (with general-purpose H/W).
 - Use it as reference for your hardware.

```
static inline float f16_to_f32(const uint32_t *input)
{
    char *iptc = (char*)input + 2;

    uint32_t half_precision = 0;
    memcpy(&half_precision, iptc + 1, 1);
    memcpy(((void*)&half_precision) + 1, iptc, 1);

    uint32_t opt = ((half_precision & 0x8000) << 16)
                  | (((half_precision & 0x7C00) >> 10) + 112) << 23
                  | ((half_precision & 0x3FF) << 13);

    float casted_opt;
    memcpy(&casted_opt, &opt, sizeof(float));
    return casted_opt;
}
```

Specification of Term Project #1

- Write your C code.
 - Map BRAM to virtual memory (use `mmap()` system call).
 - Copy DRAM's input data to BRAM.
 - Run your IP after start time is measured.
 - It will be reported as cheating if you do it before time measurement starts.
 - Modify `zynq.cpp` file only.
- Input/output
 - Input matrix : `ipt_matrix_f16`
 - Input vector : `ipt_vector_f16`
 - Output vector : `your_vector_f32`
 - After calculation is done, write your data to `your_vector_f32`

Specification of Term Project #1

- Compile and run your code
 - Compile : `$> make`
 - All warnings will be considered as error (-Werror compiler option).
 - (i.e., C++ code with warning will not be accepted.)
 - Do not modify Makefile. If you want to add compile options, ask TA for it.
 - Run : `$> ./project1`
- Use input generator
 - Usage : `$> ./input_generator > input.txt`

Review of previous labs

- It is highly recommended to review Lab materials !
 - **Lab 3:** Processing Element (PE) and BRAM model
 - This BRAM is just for simulation, you will use BRAM IP in this term project
 - **Lab 4:** PE arrays and its controller for Inner Product (FSM)
 - **Lab 6:** Block design tutorial
 - **Lab 7:** Integration of example shifter IP and its runtime debugging
 - **Lab 8:** An example full system of Inner Product unit and its customization
 - Block design is pre-described by tcl scripts, you can just start from it !
 - See how `$LAB8/hw_2` is used for simulation.
 - See how `$LAB8/hw_3` is used to generate bitstream, edit IP (myip) and run on FPGA

Demo

- Check correct output of 64x64 matrix * 64 vector output
 - This is example results with sample inputs
 - Check error rate with baseline result (32-bit * 32-bit multiplication).

```
(baseline) (CPU code) (your code)
288.495, 288.532(0.0126515%), 288.532(0.0126515%)
1696.98, 1696.95(-0.00194221%), 1696.95(-0.00194221%)
1052.98, 1052.99(0.00122884%), 1052.99(0.00122884%)
551.421, 551.488(0.0121977%), 551.488(0.0121977%)
2938.76, 2939.68(0.0312199%), 2939.68(0.0312199%)
859.224, 859.387(0.0190303%), 859.387(0.0190303%)
1911.02, 1911.03(0.000357711%), 1911.03(0.000357711%)
1090.4, 1090.28(-0.0108032%), 1090.28(-0.0108032%)
1298.14, 1298.06(-0.00676108%), 1298.06(-0.00676108%)
1026.16, 1026.25(0.0078988%), 1026.25(0.0078988%)
123.384, 123.419(0.0280729%), 123.419(0.0280729%)
2170.05, 2170.65(0.0278111%), 2170.65(0.0278111%)
287.787, 287.841(0.0185362%), 287.841(0.0185362%)
2355.97, 2356.21(0.0103937%), 2356.21(0.0103937%)
2701.55, 2700.73(-0.0305994%), 2700.73(-0.0305994%)
2221.28, 2220.89(-0.0177285%), 2220.89(-0.0177285%)
2620.62, 2620.05(-0.0215948%), 2620.05(-0.0215948%)
```